

## Module 2

Day 3

## Recap last week

- ▶ “Joined” location and property data using the url column
- ▶ Introduce `str_replace`, `str_split_fixed`, `str_trim`, and `str_to_upper` to clean variables in our data

# Goals for Today

R:

- ▶ Finish cleaning current data
- ▶ Motivate the use of user-created functions and “function safety”
- ▶ Learn how to create functions
- ▶ Create our own function to find the distance from each house to the nearest house

Economics:

- ▶ Take
- ▶ Begin exploring the relationship between distance from a metro stop and home prices

# Cleaning zip-codes

- Let's investigate our zip-code variable

```
dat_path <- ## Your file path here
joined_data <- read_csv(paste0(dat_path,
                                "joined_data.csv"))
head(joined_data$ZIP_CODE)
```

```
## # A tibble: 6 x 28
##   SALE.TYPE      SOLD.DATE PROPERTY.TYPE PRICE  BEDS BATHS SQUARE.FEET
##   <chr>          <chr>          <chr>   <dbl> <int> <dbl>         <int>
## 1 PAST SALE September-29-2017 Condo/Co-op 280000     2  1.0         1055
## 2 PAST SALE September-6-2017  Condo/Co-op 405000     2  2.0         1030
## 3 PAST SALE September-19-2017 Condo/Co-op 510000     2  2.0         1209
## 4 PAST SALE September-29-2017 Condo/Co-op 395000     2  2.0         1135
## 5 PAST SALE September-25-2017 Condo/Co-op 339900     2  1.0          930
## 6 PAST SALE September-12-2017 Condo/Co-op 415000     2  2.5         1606
## # ... with 21 more variables: LOT.SIZE <int>, YEAR.BUILT <int>,
## #   HOA.MONTH <int>, STATUS <chr>, NEXT.OPEN.HOUSE.START.TIME <chr>,
## #   NEXT.OPEN.HOUSE.END.TIME <chr>, URL <chr>, SOURCE <chr>, MLS. <chr>,
## #   FAVORITE <chr>, INTERESTED <chr>, LIST.DATE <chr>, ADDRESS <chr>,
## #   ZIP_CODE <chr>, LAT_LON <chr>, CITY_STATE <chr>, LOCATION <chr>,
## #   CITY <chr>, STATE <chr>, LATITUDE <dbl>, LONGITUDE <dbl>
```

- We can see two big problems with our current data:

## str\_sub and str\_length

- ▶ `str_length()` tells us the number of characters in a string
- ▶ `str_sub()` returns part of a text string between the start and end position provided

```
str_length("Hello")
```

```
## [1] 5
```

```
str_sub("Hello", start = 2, end = 4)
```

```
## [1] "ell"
```

## str\_sub and str\_length

- ▶ str\_length() is primarily used with other functions
- ▶ We can combine both functions in order to extract the actual zip-code from our variable

```
str_sub("'222040000", 2,  
        str_length("'222040000") - 4)
```

```
## [1] "22204"
```

- ▶ **Use dplyr and these new stringr functions to fix the rest of our zip-code data**

## Reviewing the Data

- ▶ We'll start by looking at the head of the data

```
## # A tibble: 6 x 3
##   PRICE SQUARE.FEET LOT.SIZE
##   <dbl>         <int>    <int>
## 1 280000         1055     1055
## 2 405000         1030     1030
## 3 510000         1209     1209
## 4 395000         1135     1135
## 5 339900          930      930
## 6 415000         1606     1606
```

- ▶ We have a lot of NA values. Why is that?



# Cleaning The Data

- ▶ Let's get rid of the NA values by replacing them with square footage
- ▶ We'll create a new column called Property Size that equals either lot size or square feet
  - ▶ Use an `ifelse()` statement in our `mutate`, where we test whether lot size is NA
  - ▶ If lot size is NA, use square feet. Otherwise, use lot size
- ▶ Next we will summarize our data by price per foot of property size.
  - ▶ We want to group by property type and state!

# Cleaning Code

```
joined_data <- joined_data %>%  
  mutate(PROPERTY.SIZE = ifelse(is.na(LOT.SIZE), SQUARE.FEET, LOT.SIZE))  
  
# Now we can find our average  
lot_state_data <- joined_data %>%  
  group_by(PROPERTY.TYPE, STATE) %>%  
  summarise(price = mean(PRICE/PROPERTY.SIZE, na.rm = TRUE))  
  
head(lot_state_data)
```


```
## # A tibble: 6 x 3  
## # Groups:   PROPERTY.TYPE [4]  
##           PROPERTY.TYPE STATE    price  
##           <chr> <chr>    <dbl>  
## 1           Condo/Co-op    DC 517.2998  
## 2           Condo/Co-op    MD 281.4152  
## 3           Condo/Co-op    VA 439.9952  
## 4 Multi-Family (2-4 Unit)    DC 338.9963  
## 5 Multi-Family (5+ Unit)    MD 179.8902  
## 6 Single Family Residential    DC 166.2042
```

# Visualizing The Relationship

- ▶ **Make a plot showing the average price per lot size for each property type in each state**

## Finished Plot

### Price per square foot of total property size



# Different Forms of Data

- ▶ There are two types of data that we work with:
  1. Long
  2. Wide
- ▶ Up until now we have been looking at only one format of data — long type
  - ▶ Much easier to manipulate and graph
- ▶ Wide-form data is easier to see visually
  - ▶ Generally used to display tables


# A Real Estate Example

- ▶ Let's look at a toy example with data from Redfin


```
wide_data <- read_csv(paste0(dat_path,  
                             "wide_sample.csv"))
```

- ▶ Each observation is a property at every point in time for our data
  - ▶ How is this different from an observation in our original data?
  - ▶ If this data were in long form, how many columns would there be?

# Converting Our Data



**tidyr::gather(...)**



**tidyr::spread(...)**

# gather()

- ▶ `gather(data, key, value, ...)`
  - ▶ Converts multiple columns from wide-format to long-format
  - ▶ `key` - name of the variable you are trying to create
  - ▶ `value` - variable that holds the values in the key variables
  - ▶ Must specify which columns to use for the key

```
long_data <- wide_data %>%  
  gather(key = "year", value = "price", 5:11)
```



# spread()

- ▶ `spread(data, key, value, ...)`
  - ▶ Converts two columns from long-format to wide-format
  - ▶ `key` - variable you wish to spread into multiple columns
  - ▶ `value` - variable that holds the values that will be used by your new variables

```
wide_data <- long_data %>%  
  spread(key = "year", value = "price")
```

## Other Uses

- ▶ We can use our reshaping functions in conjunction with stargazer to produce good looking tables

```
lot_state_data %>%  
  spread(key = STATE, value = price) %>%  
  stargazer(summary = FALSE,  
            header = FALSE,  
            type = "Latex")
```

## Reading In The Metro Station Data

- ▶ Use this data in conjunction with the housing data to find the distance from a metro
  - ▶ Includes latitude and longitude of each metro station
- ▶ Unfortunately it is in an xlsx file so we can't use read\_csv

```
library(xlsx)
metro_data <- read.xlsx(paste0(dat_path,
                                "Metro_lat_lon.xlsx"),
                        sheetIndex = 1)
```

- ▶ There is one big problem with finding the distance — the calculation is fairly difficult
  - ▶ We need to create a function!

# What Is A Function

- ▶ Piece of code that takes one or more inputs and returns one or more outputs
- ▶ For example:
  - ▶ `min()` takes a vector of numbers and returns the number with the lowest value
  - ▶ `read_csv()` takes a string that contains a file path and returns a `data.frame` object
- ▶ Some of inputs may have a default set value for certain arguments
  - ▶ `min(..., na.rm = FALSE)`

# Writing A Function

- ▶ Functions are extremely useful when you have to repeat a complex chunk of code
  - ▶ Can be used in loops, if else statements, and even other functions
- ▶ A user-defined function takes the following form:

```
function_name <- function(arguments){  
  ## Code that does your task  
}
```

## Simple Examples

- ▶ Here is a function for adding two numbers together:

```
add <- function(num1, num2){  
  num1 + num2  
}  
add(num1 = 2, num2 = 5)
```

```
## [1] 7
```

- ▶ Create your own function that raises a number provided by the user to a power provided by the user.

# Defaults

- ▶ Use these when one option is used much more than any of the other options
- ▶ For example `read_csv(..., col_names = T)`
  - ▶ read in the first row of the data as column names
  - ▶ may also use `FALSE` or a vector of column names we would like to assign
- ▶ Merely set one of our arguments equal to the default value

```
root <- function(base, denom = 2){  
  base**(1/denom)  
}
```

# return()

- ▶ Let's take a look the functions we created so far
  - ▶ What do we input and what do we output?
  - ▶ Is this completely clear from our code?
- ▶ By default, a function returns the last object that was created
  - ▶ Can be confusing when reading someone elses code
- ▶ `return()` says I want the output of my function to be this value
  - ▶ Only useable inside of a user-generated function



## A More Complex Example

- ▶ We can use any of the coding structures we've already learned inside of our functions
- ▶ Let's create a function that sums together the numbers in a vector

```
vector_sum <- function(numeric_vector){  
  sum <- 0  
  for(number in numeric_vector){  
    sum <- sum + number  
  }  
  return(sum)  
}  
vector_sum(c(1,1,5,12))
```

```
## [1] 19
```

## Another Complex Example

- ▶ Let's find the minimum number in a vector

```
vector_min <- function(numeric_vector){  
  m <- numeric_vector[1]  
  for(number in 2:length(numeric_vector)){  
    if(numeric_vector[number] < m){  
      m <- numeric_vector[number]  
    }  
  }  
  return(m)  
}
```

```
test_vector <- c(1, 4, 6, -8, 0, 11)  
vector_min(test_vector)
```

```
## [1] -8
```

## Another Complex Example

**Now it's your turn. Write a function called `vector_max` which takes a numeric vector and returns the maximum element in it. Make sure to test out your function.**

# Rock, Paper, Scissors

- ▶ Let's try using a more fun example now with a simple logic game

```
if(( & ) | ( & ) | ( & )){  
    print("Player 1 Wins!")  
} else if (){  
    print("Tie game!")  
} else {  
    print("Player 2 Wins!")  
}
```

**Fill in the logical arguments above. Create a function called RPS that uses the code as the body of the function and returns the results of the match.**

## Returning to Metro Data

- ▶ Remember, our goal is to calculate the nearest metro station
  - ▶ Use our longitude and latitude data combined with a little trigonometry to find the distance
  - ▶ Remember:  $a^2 + b^2 = c^2$  or  $\Delta longitude^2 + \Delta latitude^2 = distance^2$

# Calculating Distance

- Now that our data is numeric we can calculate the distance

```
test_property <- joined_data[1, ]
test_metro <- metro_data[1, ]

delta_x <- test_property$LONGITUDE - test_metro$Longitude
delta_y <- test_property$LATITUDE - test_metro$Latitude

distance <- sqrt(delta_x**2 + delta_y**2)
# One degree is equal to about 69 miles
distance/69
```

```
## [1] 0.002622848
```

## Basic Distance Function

**Write a function that takes 4 arguments: property long and lat values, and metro long and lat values. The function should then calculate the pythagorean distance between the property and the metro in miles.**

```
metro_dist <- function( , , , ){  
  
  
  
  
  
  
}
```

## Double checking the distance

- ▶ Our test property is in Arlington and our test metro is in Maryland east of D.C.
  - ▶ That's about 10 miles
  - ▶ Why might our distance be incorrect?



# geosphere

- ▶ We can't use the simple Pythagorean formula
- ▶ geosphere is a package that allows us to easily work with geographic coordinates
- ▶ `distHaversine()` calculates the shortest distance between two points
  - ▶ Points must contain longitude AND latitude
  - ▶ "As the crow flies"
  - ▶ Default output is in meters

```
# install.packages("geosphere")  
library(geosphere)  
?distHaversine()
```

## Calculating distance

```
prop_long_lat <- c(test_property$LONGITUDE,  
                  test_property$LATITUDE)  
metro_long_lat <- c(test_metro$Longitude,  
                   test_metro$Latitude)  
  
distance <- distHaversine(prop_long_lat, metro_long_lat)  
# There are 1609.344 meters in a mile  
distance/1609.344
```

```
## [1] 9.788104
```

- This looks much more realistic!

# Every Metro Stop

```
output_distances <- c()

for(metro in 1:nrow(metro_data)){
  # Pull out our longitude and latitude values for the metro station
  metro_long_lat <- c(metro_data$Longitude[metro],
                     metro_data$Latitude[metro])

  # Calculate the distance to our property
  distance <- distHaversine(p1 = prop_long_lat, # from above
                           p2 = metro_long_lat)

  # append that distance to our distances vector
  output_distances <- c(output_distances, distance)
}

output_distances
```

## Better Distance Function

Turn the code from the previous slide into a function called `prop_metros_dist` that does the following:

- ▶ Finds the distance from a property to every metro station
- ▶ Converts the distance from meters to miles
- ▶ Returns the distance (in miles) of the closest metro station

# Function safety

- ▶ Functions help use improve readability and replicability of our code, but are not fool proof
- ▶ People often input the wrong arguments or use functions in ways that weren't intended
- ▶ Let's look at `prop_metros_dist()`:
  - ▶ If a user doesn't recognize `distHaversine()` is from the `geosphere` package, they may not install and load it
  - ▶ It is our job as the programmer to address these types of concerns

# Functions for Function safety

- ▶ `require()` — works similarly to `library()`
  - ▶ Designed to be used within other functions
  - ▶ Returns `FALSE` if a package fails to load
  - ▶ Can use this with `if/else` statements
- ▶ `warning()` — prints a warning message that we specify but DOES NOT stop the program
- ▶ `stop()` — prints an error message and stops the program

# Function safety in Action

```
prop_metros_dist <- function(prop_long_lat, metro_data){  
  
  if(!require(geosphere)){ # what does require do?  
    warning("geosphere package is not installed")  
  } else{  
    output_distances <- c()  
    for(i in 1:nrow(metro_data)){  
      metro_long_lat <- c(metro_data$Longitude[i],  
                          metro_data$Latitude[i])  
      distance <- distHaversine(p1 = prop_long_lat,  
                                p2 = metro_long_lat)  
      output_distances <- c(output_distances, distance)  
    }  
  
    return(min(output_distances)/1609.344)  
  }  
}
```

# Wrapper Functions

- ▶ Need a two-number vector with longitude and latitude but we have two columns in our data
- ▶ A function that calls another function
  - ▶ Generally pre-process inputs for other functions
  - ▶ Improve readability, ease of use, and our ability to implement changes
  - ▶ `prop_metros_dist()` calls `dist_Haversine()`
- ▶ Want to write a function that turns `LATITUDE` and `LONGITUDE` into a vector



# Our Wrapper Function

```
distance_function <- function(LONG, LAT, metro_df){  
  if(!require(geosphere)){ # what does require do?  
    warning("geosphere package is not installed")  
  } else{  
    # We need to turn our LONG and LAT columns into a vector  
    prop_long_lat <- c(LONG, LAT)  
    # Now call our prop_metros_dist function from above  
    out <- prop_metros_dist(prop_long_lat,  
                           metro_df)  
    return(out)  
  }  
}  
  
distance_function(joined_data$LONGITUDE[3],  
                  joined_data$LATITUDE[3],  
                  metro_data)
```

```
## [1] 0.7918705
```

## Finding the Closest Metro Station

- **Combine our function with a loop to find the closest metro station to each property (Hint: Look at one of our previous functions)**

```
joined_data$metro_distance <- NA #initialize the column
for(){

  ## Your code here

}
```

## Graphing the Relationship